

01 | What is Data Science

Ivan Corneillet

Data Scientist

Learning Objectives

After this lesson, you should be able to:

- Describe the components of a successful learning environment
- Define what is data science and who data scientists are; define the data science workflow
- Setup your development environment and practice the different workflows we will use in this course

A black circle containing the white text "DS".

DS

Setting You Up for Success

Meet Your Team

- Ivan Corneillet, Lead Instructor



- TBD, Associate Instructor

- Matt Jones, Course Producer



Course Logistics


- Lead Instructor
 - Ivan Corneillet
- Associate Instructor
 - TBD
- Course Producer
 - Matt Jones
- Class
 - August 29 – November 2, Tuesdays and Thursdays, 6:30PM – 9:30PM
 - Classroom 4
- Slack
 - <https://ds-sf-38.slack.com>
- GitHub
 - <https://github.com/ga-students/DS-SF-38>
- Exit Tickets
 - <http://tiny.cc/ds-sf-38>

What skills will I learn in this class?

What is Data Science (session 1)	Research Design (session 1)	Python (session 2)	<i>pandas</i> (session 3)	Databases and Scrapping (session 4)	Exploratory Data Analysis (session 5)
<i>k</i> -Nearest Neighbors (session 6)	Model Fit (session 6)	Data Wrangling and Exploratory Data Analysis Challenge (session 7)	Linear Regression (sessions 8–10)	Regularization (session 11)	Logistic Regression (session 12)
Machine Learning Modeling Challenge (session 13)	Trees (session 14)	Review (session 16)	Machine Learning Modeling Challenge, Take 2 (session 17)	Natural Language Processing (session 18)	Time Series (session 19)

Top algorithms and methods used by data scientists


(<http://www.kdnuggets.com/2016/09/poll-algorithms-used-data-scientists.html>)



Gladwell's 10,000 Hour Rule

(<http://www.wisdomgroup.com/blog/10000-hours-of-practice>)

- ▶ “Greatness requires enormous time”
 - ▶ It takes roughly ten thousand hours of practice to achieve mastery in a field



How will I apply and reinforce these new skills?

Assignments <i>(ungraded)</i>	Take-home assignments		
Unit Project <i>(graded)</i>	Research Design <i>(session 5)</i>	Exploratory Data Analysis <i>(session 9)</i>	Machine Learning Modeling and Executive Summary <i>(session 14)</i>
Challenges <i>(ungraded)</i>	Data Wrangling and Exploratory Data Analysis <i>(session 7)</i>	Machine Learning Modeling <i>(session 13)</i>	Machine Learning Modeling <i>(session 17)</i>
Final Project <i>(graded)</i>	Lightning Pitch <i>(session 10)</i>	Research Design, Exploratory Data Analysis, and Intermediate Presentation <i>(session 15)</i>	Machine Learning Modeling and Final Presentation <i>(session 20)</i>

Typical Class

- *Pre-readings (usually optional)*
- Objectives
- Announcements
- Previous class review
- Series alternating between:
 - Lectures
 - (deck, whiteboard, codealongs, and demos)
 - Activities
 - (cold-calling, individual and group exercises, and codealongs)
- Class review
- Exit tickets
- *Post-readings (usually optional)*



DS

What is Data Science? Who are Data Scientists?

Harvard Business Review | “Data Scientists: The Sexiest Job of the 21st Century” (2012)


(<https://hbr.org/2012/10/data-scientist-the-sexiest-job-of-the-21st-century>)



Source: Harvard Business Review

Wired | “The End of Code” | “Soon We Won't Program Computers. We'll Train Them Like Dogs” (2016)

(<http://www.wired.com/2016/05/the-end-of-code>)



Source: Wired

Data science is everywhere

 **FiveThirtyEight**

NETFLIX



Walmart 



amazon 

Google

TARGET



U B E R



Linked in

Common questions asked in data science

How much? How many?

- What will the temperature be next Tuesday?
- What will my fourth quarter sales in France be?
- How many kilowatts will be demanded from my wind farm 30 minutes from now?
- How many new followers will I get next week?

Regression

- Predict a continuous outcome
 - k -Nearest Neighbors
 - Linear Regression
 - Trees

Common questions asked in data science (cont.)

Is this A, B or C?

- Will this customer default on their loan?
- Is this an image of a man, a cat, or a dog?
- Will this customer click on the advertisement?
- Which team will win the championship?
- Is this mole malignant or benign?

Classification

- Predict a discrete outcome
 - k -Nearest Neighbors
 - Logistic Regression
 - Trees

Common questions asked in data science (cont.)


How is this Data Organized?

- What are the different types of coffee drinkers?
- Which viewers like the same kind of movies?
- What kinds of car models does GM produce?
- Are there common clusters of cable channels that customers tend to purchase together?
- What is a natural way to break these documents into five topics?

Clustering

- What are the “categories” within the data?

Data science involves a variety of skillsets



Source: Data Science for the C-suite

Data scientists in ≤ 140 characters

The image displays four tweets, each enclosed in a colored border, arranged in a square layout. In the center of the square is a Twitter logo with the letters A, B, C, and D positioned around it. The tweets are as follows:

- Zvi (@nivertech)**: "Data Scientist" is a Data Analyst who lives in California.
RETWEETS: 162, LIKES: 82
5:55 PM - 14 Mar 2012
- Josh Wills (@josh_wills)**: Data Scientist (n.): Person who is better at statistics than any software engineer and better at software engineering than any statistician.
RETWEETS: 1,339, LIKES: 799
9:55 AM - 3 May 2012
- Harvard IACS (@Harvard_IACS)**: Data Scientist: Someone better at statistics than a software engineer, and better at software engineering than a statistician?
#datastorm14
RETWEETS: 43, LIKES: 15
6:42 AM - 24 Jan 2014
- Javier Nogales (@fnogales)**: Data Scientist (2/2): person who is worse at statistics than any statistician and worse at software engineering than any software engineer
RETWEET: 1, LIKES: 5
6:08 AM - 27 Jan 2014

Source: Twitter

Wired's "Soon We Won't Program Computers. We'll Train Them Like Dogs" (2016) (cont.)

Behaviorism/Behavioral Psychology

- Brain as a black box
 - Stimulus and response, feedback and reinforcements
 - "ring bell, dog salivates"

Cognitive Psychology

- Brain more like a computer
 - Thoughts as programs
 - Absorb, process, and act upon information

Wired's "Soon We Won't Program Computers. We'll Train Them Like Dogs" (2016) (cont.)

Machine Learning

- Humans *train* computers
 - Keep showing cats to a computer and eventually it will *learn* to recognize cats (<https://www.wired.com/2012/06/google-x-neural-network>)
 - No symbols, no rules; instead an unparsable machine learning

Traditional Programming

- Humans *write code* (as explicit step-by-step-instructions) for computers to follow
 - Rule-based determinism
 - "Write enough rules and eventually, we'd create a system sophisticated enough to understand the world"
 - For years, Google Search relied mostly on these human-written rules (<https://www.wired.com/2016/02/ai-is-changing-the-technology-behind-google-searches>)

Wired's "Soon We Won't Program Computers. We'll Train Them Like Dogs" (2016) (cont.)

Age of Entanglement

- Outside-in view of how machine work
 - "Code doesn't just determine behavior, behavior also determine code"

Age of Enlightenment

- Inside-out view of how machine work
 - "First, we write the code, then the machine expresses it"

In this course we will model the stimuli as a matrix \mathbf{X} (the **feature matrix**); the response is modelled as a vector \mathbf{y} (the **response vector**). We will use these key data structures as inputs to our machine learning algorithms

Feature Matrix \mathbf{X}

Stimulus/feedback
"ring bell"

	col0	col1	col2	col3
row0				
row1				
row2				
row3				

Response Vector \mathbf{y}

Response/reinforcements
"dog salivates"

	col
row0	
row1	
row2	
row3	




DS

Data Science Workflow



Data Science Workflow

(a.k.a., Cross-Industry Standard Process for Data Mining, or CRISP-DM)




Some resources to follow along the class (or afterwards...) (*will reference either pre-class reading and/or post-class reading materials; optional; not required for the course*)

- Data Science for Business (by Provost and Fawcett) ([link](#)) (General Assembly holds several copies in its library)



- An Introduction to Statistical Learning: with Applications in R (by James et al.) (e-book available free-of-charge [here](#))

- For a more advanced treatment of these topics, check out The Elements of Statistical Learning: Data Mining, Inference, and Prediction (by Hastie et al.). (e-book also available free-of-charge [here](#))



Slides © 2017 Ivan Corneillet Where Applicable
Do Not Reproduce Without Permission